

# Linear classifiers

Gosia Migut

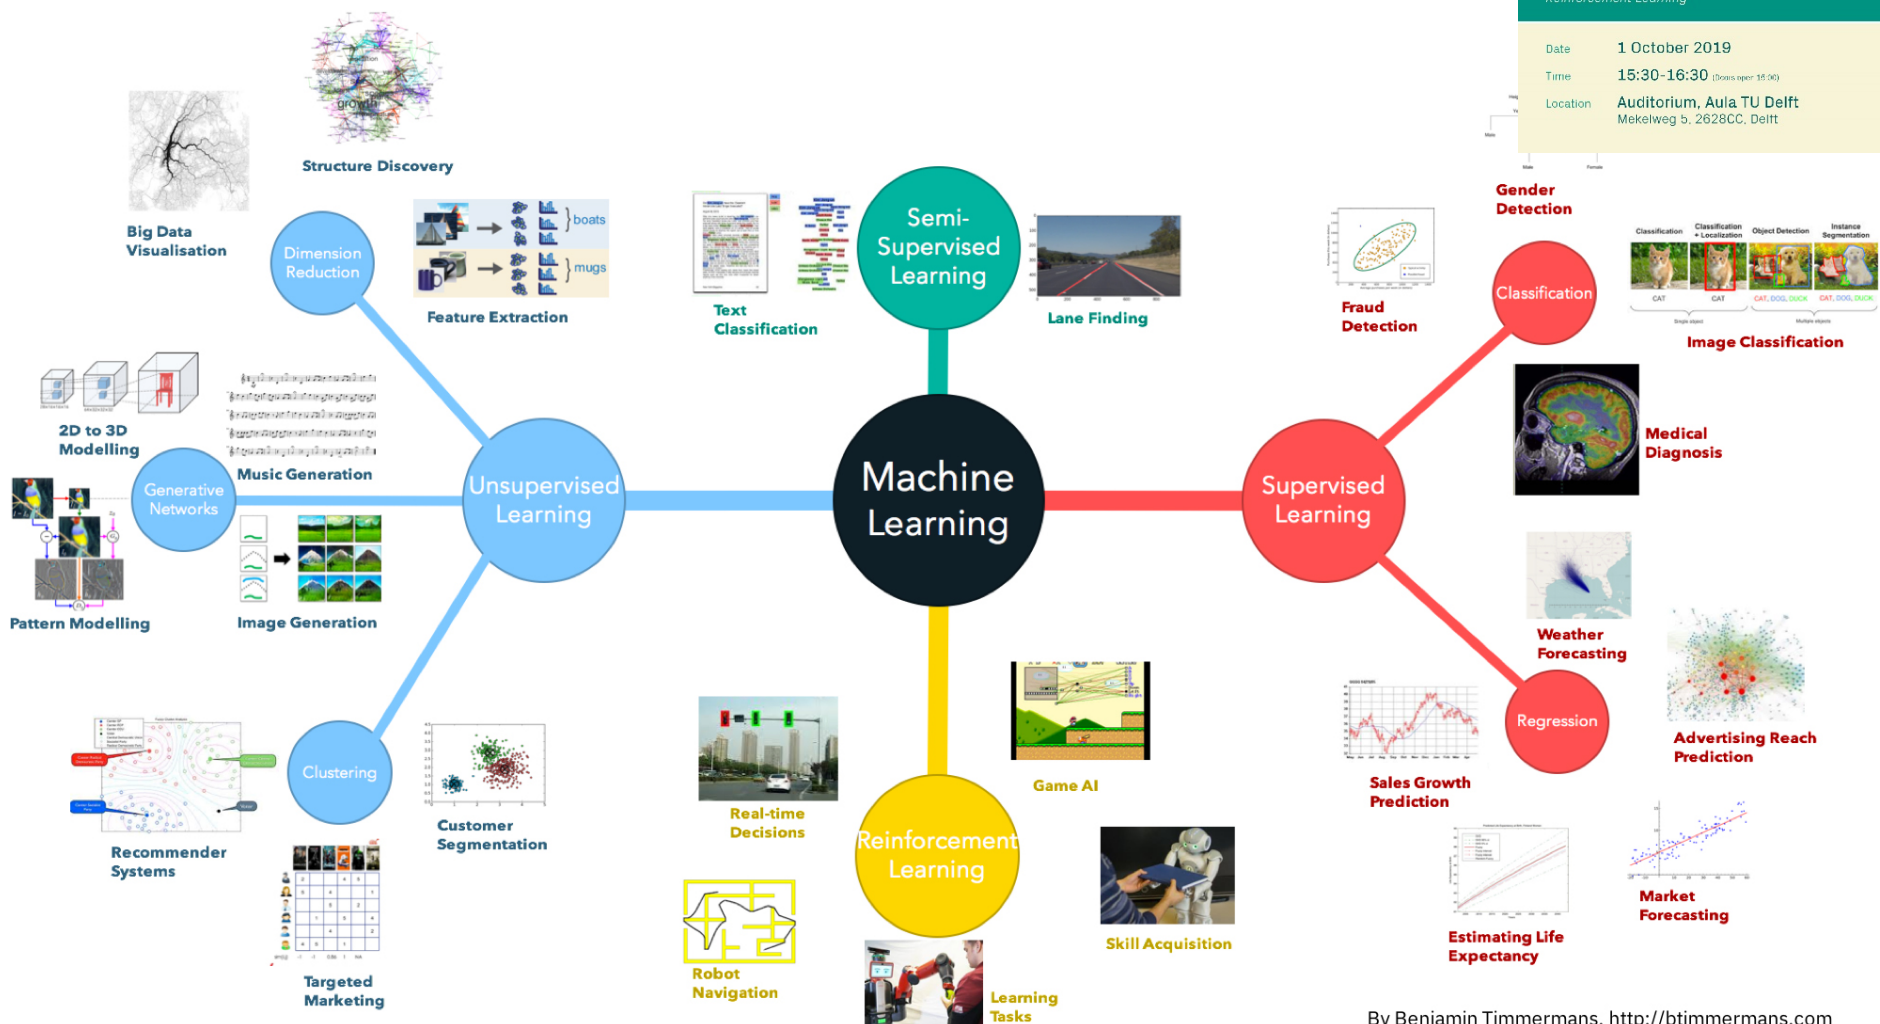
# Admin stuff

- Feedback (thank you, also student panel!)
  - I'm staying!
  - 2nd year bachelor CSE course
  - Digital practice exam in week 5
  - No labs answers, TA's available
  - More examples with the formulas
- Lab week 4 more challenging!
  - Includes material from Tuesday and Friday lecture
  - Notation corresponds with the reading material

# I “owe” you

- Will be fixed this week:
  - Pseudo code Parzen width parameter optimization
  - Solution last exercise Naive Bayes

# Machine Learning overview



**Pieter Abbeel**  
*Deep Learning to Learn*  
 Keynote about the State of the Art in Reinforcement Learning

Date 1 October 2019  
 Time 15:30-16:30 (Doors open 15:00)  
 Location Auditorium, Aula TU Delft  
 Mekelweg 5, 2628CC, Delft

# Generative vs discriminative models

- A **generative** model explicitly models the joint probability distribution  $p(x|y)$  and then uses Bayes rule to compute posterior probabilities  $p(y|x)$ 
  - Parametric density estimation: eg. Nearest mean, LDA, QDA
  - Non-parametric density estimation: eg. k-nn, Parzen, Naive Bayes
- A **discriminative** model directly models  $p(y|x)$  from the training examples.
  - Linear: eg. logistic regression, svm
  - Non-linear: eg. decision trees, multi-layer perceptron

# Learning objectives of this lecture

- After this lecture you will be able to explain:
  - what the general idea of linear classification is
  - what  $w^T x$  means
  - What a cost function is
  - the gradient descent algorithm
  - how to optimize a cost function using gradient descent
  - what the difference between gradient descent and stochastic gradient descent is.

# Reading of this week

- CS229 Lecture notes by Andrew Ng (Stanford University):
  - Supervised learning p.1-2
  - Part I Linear regression p. 3-4
    - 1. LMS algorithm p. 4-7
    - 3. Probabilistic interpretation p. 11-13
  - Part II Classification and logistic regression p. 16-19

<http://cs229.stanford.edu/notes/cs229-notes1.pdf>
- Lab of week 4 is consistent with the notation of the reading

# Linear classifiers



# Note on the notation

- Parameters notation
  - In the lecture we use  $\mathbf{w}$
  - In the reading and lab  $\theta$  is used

# Linear classifier

- Linear classifier has a **linear decision boundary**.
- Decision boundary of a linear classifier for 2 dimensions is a line:


$$w_1x_1 + w_2x_2 + w_0 = 0$$

- A hyperplane is a generalization of a straight line to  $> 2$  dimensions.
- A hyperplane contains all the points in a  $d$  dimensional space satisfying the following equation

$$w^T x + w_0 = 0$$

# Linear classifier: terminology

$$h(x) = w^T x + w_0$$

- The slope of the hyperplane is determined by the **parameter (weight) vector**  $w = (w_1, \dots, w_d)$  . 
- The location (intercept) is determined by **bias**  $w_0$  .
- The function of the input  $h(x)$  is a **linear combination** of the parameters  $w$  .

# Linear classifier

- Given the linear classifier:


$$h(x) = w^T x + w_0$$




- Classify  $x$  to  $\begin{cases} y_1 & \text{if } w^T x + w_0 \geq 0 \\ y_0 & \text{if } w^T x + w_0 < 0 \end{cases}$

# What does $w^T x$ mean?

- Assume I have a 2-dimensional  $w$ :



# What does $w^T x + w_0$ mean?



$w_0 = ?$

# Incorporate the bias term

- Quite often you see

$$h(x) = w^T x = 0$$

- Instead of

$$h(x) = w^T x + w_0 = 0$$

$$h(x) = w_1 x_1 + w_2 x_2 + \dots + w_d x_d + w_0$$


- No problem if we redefine the feature vector:

$$\tilde{x} = \begin{bmatrix} x \\ 1 \end{bmatrix} \Rightarrow x_0 = 1$$

- $h(x) = w_0 x_0 + w_1 x_1 + w_2 x_2 + \dots + w_d x_d = \sum_{i=0}^d x_i w_i = w^T x$



# Linear classifier



- Classifier is a linear function of the features
- The classification depends if the weighted sum of the features is above or below 0



# Linear classifier


- The goal of the learning process is to come up with a “good” weight vector  $w$
- The learning process will use examples to guide the search of a “good”  $w$
- Different notions of “goodness” exist, which yield different learning algorithms

# Define a “goodness”/error measure

- Cost function
  - Measure of performance => single real number
  - Should be optimized
  - Yields different learning algorithms
  - Eg. Log-likelihood (Naive Bayes)

# Cost function

# Consider a regression problem...



Linear Regression?

I covered that last year.


Wake me up when we get to  
Support Vector Machines!

Noah Mackey

# Univariate linear regression

- Training data: observations paired with outcomes (real number)
- Observations have features (predictors, typically also real numbers)
- The model is a regression line  $y = w_1x + w_0$  which best fits the observations:
  - $w_1$  is the slope
  - $w_0$  is the intercept
  - This model has two parameters (or weights)
  - One feature =  $x$
  - Example:
    - $x$  = size of property
    - $y$  = price of property

# Linear regression




# Multivariate linear regression

- More generally  $y = w_0 + \sum_{i=1}^d w_i x_i$ , where
  - $y$  is outcome
  - $w_0$  is intercept
  - $x_1, \dots, x_d$  is feature vector and
  - $w_1, \dots, w_d$  parameter/weight vector
- Get rid of bias:  $\sum_{i=1}^d x_i w_i = w^T x$

# Cost function for linear regression

- Minimize sum squared error over N training examples






# Cost function intuition

- Hypothesis:  $h(w) = w_1 x_1 + w_0$
- Parameters:  $w_0$  and  $w_1$
- Cost function:  $J(w_0, w_1) = \frac{1}{2n} \sum_{i=1} (h(x)^{(i)} - y^{(i)})^2$
- Goal: minimize  $J(w_0, w_1)$   
 $w_0, w_1$

# Cost function intuition

- Cost function  $J(w_1)$  against  $w_1$



# Cost function optimization

- Solution: Set the derivative to 0, and solve:

$$\frac{\partial J(w)}{\partial w} = 0$$

(typically hard/impossible to do)

- Solution: Follow the derivatives (gradient) until you hit a (local) minimum.
  - What is gradient descent?
  - What is stochastic gradient descent?

# Gradient descent

# Gradient descent algorithm

- Goal: minimize  $J(w_0, w_1)$   
 $w_0, w_1$
- Outline:
  - Start with some  $w_0, w_1$  (eg.  $w_0 = 5, w_1 = 0.167$ )
  - Keep changing  $w_0, w_1$  to reduce  $J(w_0, w_1)$
- Repeat until convergence {

$$w_j := w_j - \alpha \frac{\partial J(w_0, w_1)}{\partial w_j}$$


}

What does it all mean?

# Gradient vector

$$w_j := w_j - \alpha \frac{\partial J(w_0, w_1)}{\partial w_j}$$


- Gradient vector has as coordinates the partial derivatives of a function
- $\alpha$  is learning rate = speed of descent



# Learning rate


$$w_j := w_j - \alpha \frac{\partial J(w)}{\partial w_j}$$

- If  $\alpha$  is too small gradient descent can be slow
- If  $\alpha$  is too large, gradient descent can overshoot the minimum (it may fail to converge)




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )



$$J(w_0, w_1)$$


(function of the parameters  $w_0, w_1$ )






$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$h(x)$$

(for fixed  $w_0, w_1$  this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )





# Gradient descent



# Gradient descent for univariate linear regression

- Hypothesis:  $h(w) = w_1 x_1 + w_0$
- Parameters:  $w_0$  and  $w_1$
- Cost function:  $J(w_0, w_1) = \frac{1}{2n} \sum_{i=1} (h(x)^{(i)} - y^{(i)})^2$
- Goal: minimize  $J(w_0, w_1)$   
 $w_0, w_1$

# Gradient descent for univariate linear regression

- $w_0 := w_0 - \alpha \frac{\partial(\frac{1}{2n} \sum_{i=1} ((w_1 x_1 + w_0)^{(i)} - y^{(i)})^2)}{\partial w_0}$
- $w_0 := w_0 - \alpha \frac{1}{n} \sum_{i=1} ((w_1 x_1 + w_0)^{(i)} - y^{(i)})$
- $w_1 := w_1 - \alpha \frac{\partial(\frac{1}{2n} \sum_{i=1} ((w_1 x_1 + w_0)^{(i)} - y^{(i)})^2)}{\partial w_1}$
- $w_1 := w_1 - \alpha \frac{1}{n} \sum_{i=1} ((w_1 x_1 + w_0)^{(i)} - y^{(i)}) x_1^{(i)}$

# General idea of gradient descent

- A gradient is a slope of a function
- That is, a set of partial derivatives, one for each dimension (parameter)

$$w_j := w_j - \alpha \frac{\partial J(w)}{\partial w_j}$$

- By following the gradient of a function we can descend to the minimum
- $\alpha$  is a learning rate and controls the speed of descent

# Stochastic gradient descent

- We could compute the gradient of cost function for the full dataset before each update
- Instead
  - Compute the gradient of the cost function for a single example
  - Update the weight
  - Move on to the next example

# Logistic regression

# Logistic regression: a taste

- So that you can start with the lab
- More details on Friday


# Logistic regression

- Let's change the form of linear hypotheses

$$h(x) = w^T x \text{ to satisfy } 0 \leq h(x) \leq 1$$

$$g(z) = \frac{1}{1+e^{-z}}$$


- Let's plug  $w^T x$  into the logistic function
- $z = w^T x$
- $h(x) = g(w^T x)$





# Logistic regression properties

- $h(x) = \frac{1}{1+e^{(-w^T x)}}$



- $h(x)$  will give us the probability that our output is 1
- $g(z) \rightarrow 1$  as  $z \rightarrow \infty$
- $g(z) \rightarrow 0$  as  $z \rightarrow -\infty$
- Why are these properties convenient to model a probability?

# On Friday

- More on Logistic regression and it's cost function
- Example how to calculate 1 step of gradient descent for logistic regression
- Support Vector Machine and it's cost function
- Multi-class classification
- Bias vs. variance